# RAG Overview

Retrieval-augmented generation combines a search step with language-model generation. A typical pipeline ingests source documents, extracts text, splits it into overlapping chunks, creates an embedding for each chunk, stores those vectors, retrieves the most relevant chunks for a question, and gives the retrieved context to a language model. Retrieval grounds the answer in the source corpus and can reduce unsupported claims.