Vector Search Guide

An embedding is a numeric representation of text in which semantically similar texts tend to be close together. Cosine similarity measures the angle between two vectors; a larger cosine similarity generally means the texts are more semantically alike. A vector database indexes embeddings and returns nearest neighbors efficiently. Metadata filters can restrict results to a file, page range, or other attribute.